

Lecture 18: Translation Lookaside Buffer (TLB)


Vivek Kumar

Computer Science and Engineering

IIIT Delhi


vivekk@iiitd.ac.in

Last Lecture



- Paging divides the available memory into small and fixed size blocks (4KB in Linux)
- Each VA or PA belongs to a VP or PP, respectively, and the corresponding numbers are VPN and PPN, respectively
- Linux uses a mechanism known as **Page Fault** for lazy page allocation

How to Implement Simple Paging?



Today's Class

- TLB
- Locality

Mapping Virtual to Physical Address

● Approaches

1. Virtual address is same as the physical address
2. Virtual address is not the same as physical address
 - a) Using base & bound for VA to PA mapping
 - b) Using segmentation for VA to PA mapping
 - c) Paging
 - a) Using Page Table (PT) in RAM and having a base register for PT
 - b) **Using Page Table (PT) in RAM and having a base register for PT as well as a TLB cache**

Page Table Entry

- Recall, a page table is a mapping between VPN to PPN
 - What kind of data structures we can use for implementing a page table?
 - Arrays?
 - As of now, we have seen an array based implementation of Page Table
 - VPN will be the index in an array of PPNs
 - Hash table?
 - VPN will be the key with the value as PPN
 - Overheads?
 - Space?
 - Time?

Overheads with the Page Table

● Time overhead


Discussion for today

- Accessing any virtual address requires extra memory access
 - Step-1 is to fetch the page table base address from register (PTBR)
 - Single cycle access
 - Step-2 is to index into the PTE and fetch the PPN
 - 10s to 100s of cycles depending on whether PT is in cache or RAM


● Space overhead

- Next lecture

Detour: Memory Hierarchy



Detour: How Bad is Latency as we go Down?




- Another analogy
 - Normalizing with L1 latency, and assuming one seconds is equal to 4 cycles
 - L1 = one second
 - L2 = 7.5 seconds
 - Main memory = 2.5 minutes
 - Hard drive = in several days!


Animation source: <https://overbyte.com.au/misc/Lesson3/CacheFun.html>

Detour: General Cache Concepts


- **Cache:** A smaller, faster storage device that acts as a staging area for a subset of the data in a larger, slower device



Detour: General Cache Concepts



Detour: General Cache Concepts



Data needed: 12

Data is not in cache:
Miss!


*Data is fetched from
memory*

Data is stored in cache
By evicting some old data

So, How to Solve Slow Access of PT?

- Provide a hardware cache per CPU for saving frequently accessed PTE entries (different than the CPU cache L1, L2, etc.)
 - Translation Lookaside Buffer (TLB) cache that is a part of MMU
 - Typically accessed in a single cycle
- Save most recent VPN->PPN mappings on TLB (plus some more info, but not PT entry)



TLB Hit and Miss



- **TLB hit**
 - If entry already in TLB then have the physical address without reading the page table
 - VPN(4)->PPN(10) is already on TLB. Hence, no need to access PT
- **TLB miss**
 - VPN(10)->PPN(9) mapping is not found on TBL
 - Locate PT using PTBR
 - Access page table to find the mapping and save the mapping on cache by evicting an earlier entry

Locality Helps in Caching

- Principal of Locality
 - Empirical observation: Processors tend to access same set or nearby memory locations repetitively over a short period of time
- Temporal locality:
 - Recently referenced items are likely to be referenced again in the near future
- Spatial locality:
 - Items with nearby addresses tend to be referenced close together in time



Locality Helps in Caching

- Caching helps only if application supports **locality**
 - Let us understand the concept of **locality** using a simple application called as **iterative averaging**
 - Iterative averaging is the process of updating an array to so that each index becomes the average of the indices one before and one after it
 - After repeating this for many iterations, the array may converge to one set of numbers
 - Repetitive pattern in several scientific applications

Iterative Averaging

```
//all elements zeroed
float A[SIZE+2], B[SIZE+2];

void iterative_averaging() {
    A[SIZE+1] = B[SIZE+1] = 1;
    for (int iter=0; iter<ITERATIONS; iter++) {
        for (int j=1; j<=SIZE; j++) {
            B[j] = (A[j-1] + A[j+1])/2.0;
        }
        double* temp = B;
        B = A;
        A = temp;
    }
}
```

https://classes.engineering.wustl.edu/cse231/core/index.php/Iterative_Averaging

- For SIZE=9, the array A will eventually converge with values as 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9

Spot the Locality


```
//all elements zeroed
float A[SIZE+2], B[SIZE+2];

void iterative_averaging() {
    A[SIZE+1] = B[SIZE+1] = 1;
    for (int iter=0; iter<ITERATIONS; iter++) {
        for (int j=1; j<=SIZE; j++) {
            B[j] = (A[j-1] + A[j+1])/2.0;
        }
        double* temp = B;
        B = A;
        A = temp;
    }
}
```

https://classes.engineering.wustl.edu/cse231/core/index.php/Iterative_Averaging

- Let us only consider the locality in array accesses
- Spatial locality
 - Inner for-loop
 - Reference array elements in succession (stride-1 reference pattern)
- Temporal locality
 - Outer for-loop
 - Accessing each arrays elements repeatedly


Spatial Locality



- Assumptions
 - The only memory access required are for accessing elements of arrays A and B
 - SIZE=6
 - Page size (PS) of 16 bytes
 - A requires 2 pages
 - B requires 2 pages
 - **TLB can hold 4 entries**
- Inside each iteration of outer loop
 - B is accessed **6** times and A is accessed **12** times (spatial locality)
 - Total memory accesses = 12+6=**18**
 - TBL misses
 - **2+2 = 4**
 - TLB hits
 - **10+4 = 14**

```
for (int j=1; j<=SIZE; j++) {
    B[j] = (A[j-1] + A[j+1])/2.0;
}
```

Temporal Locality




- Same indices are repeatedly accessed across the iterations of the outer for-loop
 - **Temporal locality**
- Total TLB misses for N number of iterations of outer for-loop?
 - **4 only!**
- Total page table accesses with TLB?
 - **4 only!**
- Total page table access without TLB?
 - **18N**

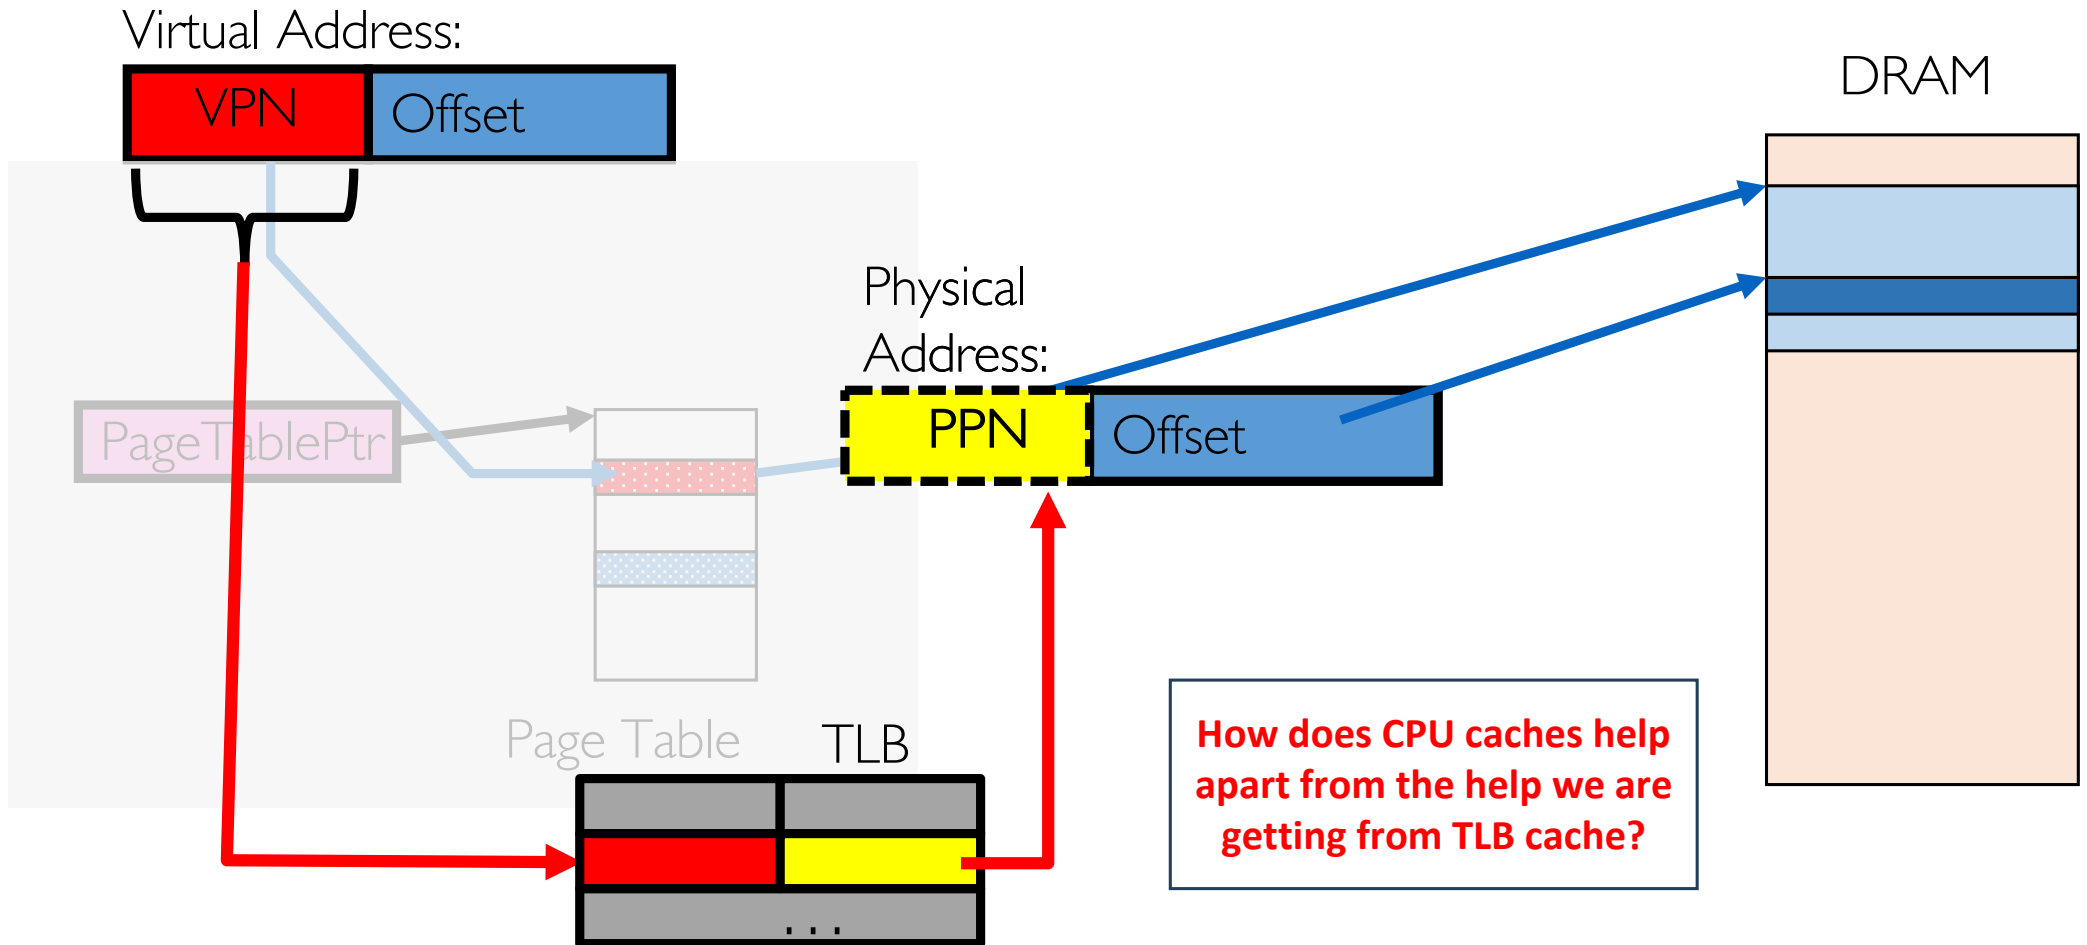
What to do at Context Switch?

- Need to do something, since TLBs map virtual addresses to physical addresses
 - Address Space just changed, so TLB entries no longer valid!
- Options?
 - Invalidate (“Flush”) TLB: simple but might be expensive
 - What if switching frequently between processes?
 - Include ProcessID in TLB
 - This is an architectural solution: needs hardware

Putting Everything Together: Address Translation



Putting Everything Together: TLB



Next Lecture

- Space overheads with the page table